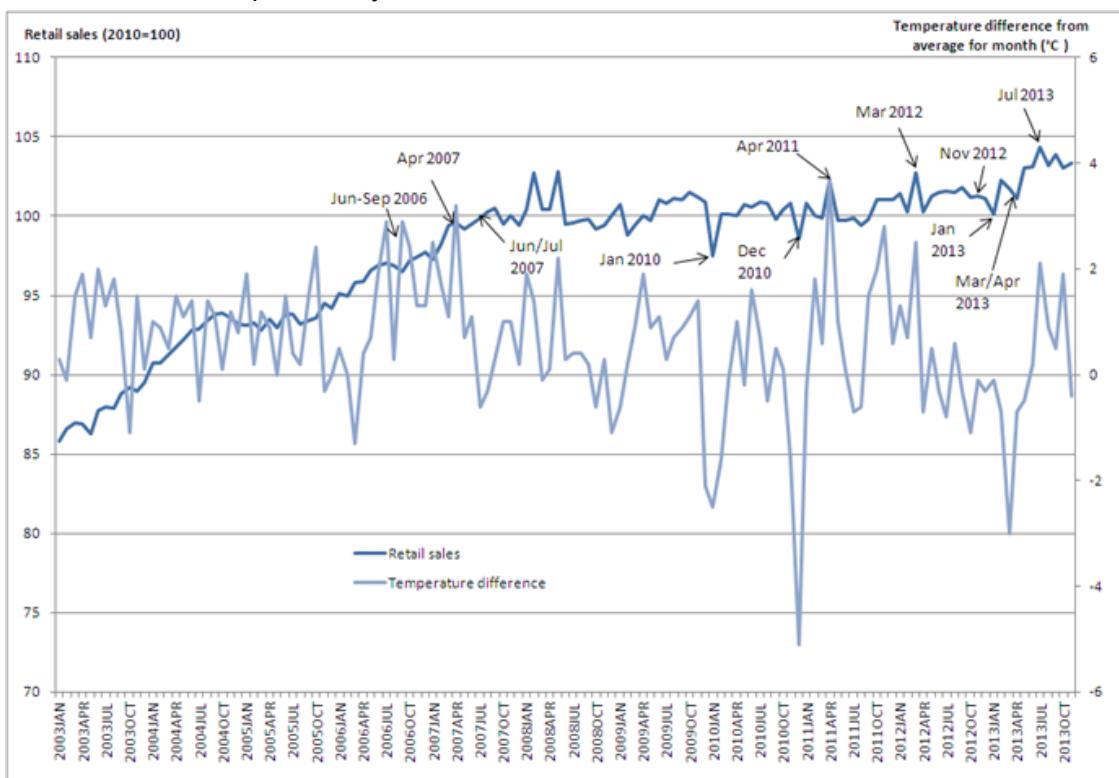


Report:

Nir Cafri

Introduction:

- product engagement increases sales - consumers pass through stages of increasing commitment (identification, enthusiasm, attention, absorption and interaction) before achieving full engagement with an offering.¹
- War decreases buying power²
- Food recalls affect purchases (depends on country and recall strategy)³
- Weather affects retail sales: “extra sales growth (%) when the temperature is 1 degree Celsius higher , c.q. there is 10 hours more precipitation in that month, compared to the same month of the previous year.”^{4,5}



¹ Woodall, T., Rosborough, J., & Harvey, J. (2018). Proposal, project, practice, pause: Developing a framework for evaluating smart domestic product engagement. *AMS Review*, 8(1), 58-74.

² <https://www.investopedia.com/terms/p/purchasingpower.asp>

³ Hu, H., Djebarni, R., Zhao, X., Xiao, L., & Flynn, B. (2017). Effect of different food recall strategies on consumers' reaction to different recall norms: A comparative study. *Industrial Management & Data Systems*.

⁴ <https://www.weatherads.io/blog/the-impact-of-weather-on-retail-sector-in-the-uk>

⁵ Fris, P. (1997). Disappearing clouds: Weather influences on retail sales. In *New Operational Approaches for Financial Modelling* (pp. 195-206). Physica, Heidelberg.

CV data clustering:

Data preprocessing:

Only actions above score threshold were taken.

Score threshold was set to $2/(number\ of\ labels)$ in order to get significant actions only.

Actions were split into 2 groups: actions which engage with the product:

```
'hold_product': 1, 'pick_up': 2, 'put_to_cart': 3, 'touch_product': 4
```

And actions which do not engage with the product:

```
'stand': 7, 'walk': 8, 'put_back': 6
```

Stop at category was dropped from all actions due to redundancy(high probability and low informativity).

Division to groups was done due to inconsistent CV actions above score threshold.

Matrices chosen are:

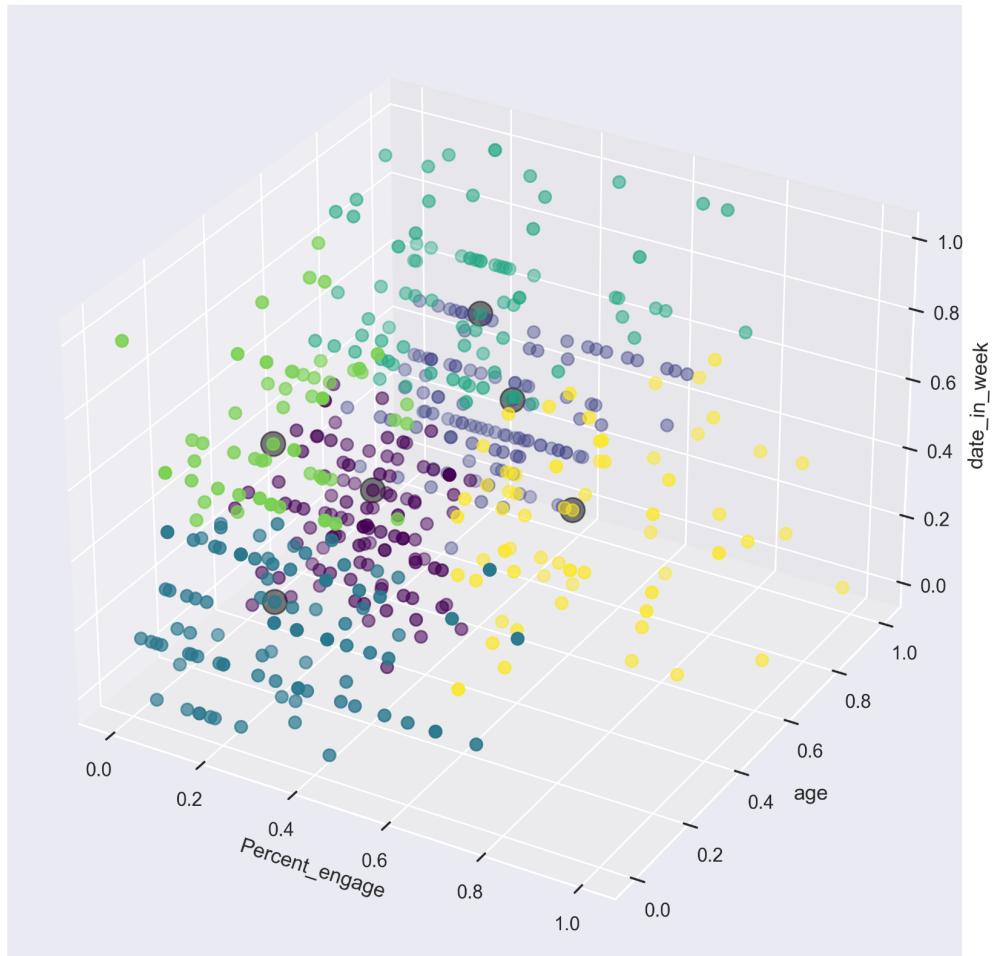
- Percent actions of engagement with the product out of all actions.
- Day in week
- Age of customer

Additional information extracted from data and used for combined model:

- Mean location (based on x,y) of actions which engage with the product
- Mean location (based on x,y) of actions which don't engage with the product
- Date

Clustering results:

Clustering method K-medoids was used due to high modulation of the algorithm. I used built centroids initialization and multiple initialization epochs due to datas' high heterogeneity. This method has a high probability of clustering which would be consistent.



Results show a few groups of low engagement with the product. Those groups were consistent with 4-7 clustering groups. And one group in yellow shows a high engagement percentage. These active users are especially engaged with the product and are mainly an older group who shops later in the week.

BI Modeling: Data preprocessing:

Excels' data was converted with external features added per branch for each row to Pandas Dataframe. A run of all the excels took a few hours. Loading bar was added to monitor progress:

Pandas dataframe was saved to .csv in saved_data folder to allow faster revival of data.

All concatenation data was saved in “saved data” in a .csv file for faster load.

Labels - Pidyon column

Features matrix -

All string data was factorized to categorical integer per unique value of column.

Index and BarCode column was dropped.

$$x_{new} = \frac{x - x_{min}}{x_{max} - x_{min}}$$

All data values were normalized:

Features added:


- Weather. Based on location of store. Including temperature, wind and rain. Weather was used with geopy. To get location I used meteosat. Better to use GoogleMaps but it requires payment(or at least registering with a credit card).
- Google trends - War and recall. Change in interest of the last 2 days compared to the last two weeks.

Modeling:

Since labeling is continuous I used regression. For tabular data I use tabnet which is based on pytorch.

Tabnet⁶ uses DL to optimize short term forecasting ⁷

Tabnet model:



In addition in order to avoid overfitting I used k fold cross validation (CV).

Hyperparameters were optimized with grid search & random tuning using Ray tuning:

⁶ Arik, S. O., & Pfister, T. (2021). Tabnet: Attentive interpretable tabular learning. In AAAI (Vol. 35, pp. 6679-6687).

⁷ Borghini, E., & Giannetti, C. (2021). Short Term Load Forecasting Using TabNet: A Comparative Study with Traditional State-of-the-Art Regression Models. *Engineering Proceedings*, 5(1), 6.

Example:

Number of trials: 22/45 (1 RUNNING, 21 TERMINATED)											
Trial name	status	loc	lr	momentum	n_d	n_steps	weight_decay	iter	total time (s)	rmse	
train_model_4a7e5_00021	RUNNING	127.0.0.1:4960	0.01	0.060939	24	8	0.000358961				
train_model_4a7e5_00000	TERMINATED	127.0.0.1:4960	0.01	0.051836	8	6	0.00023708	1	42.4906	365.64	
train_model_4a7e5_00001	TERMINATED	127.0.0.1:4960	0.001	0.0893637	8	7	9.90463e-05	1	46.6516	383.216	
train_model_4a7e5_00002	TERMINATED	127.0.0.1:4960	0.0001	0.0350622	8	4	0.000750318	1	32.5725	384.124	
train_model_4a7e5_00003	TERMINATED	127.0.0.1:4960	0.01	0.054822	24	9	3.84049e-05	1	51.1531	353.599	
train_model_4a7e5_00004	TERMINATED	127.0.0.1:4960	0.001	0.0274308	24	5	0.000103287	1	37.6657	383.249	
train_model_4a7e5_00005	TERMINATED	127.0.0.1:4960	0.0001	0.0793574	24	9	0.000139107	1	11.3556	384.656	
train_model_4a7e5_00006	TERMINATED	127.0.0.1:4960	0.01	0.0352161	32	1	0.000985244	1	13.1897	365.427	
train_model_4a7e5_00007	TERMINATED	127.0.0.1:4960	0.001	0.0408989	32	1	0.000940535	1	13.7741	383.61	
train_model_4a7e5_00008	TERMINATED	127.0.0.1:4960	0.0001	0.0158825	32	7	0.000176997	1	46.2798	384.142	
train_model_4a7e5_00009	TERMINATED	127.0.0.1:4960	0.01	0.0905694	8	5	7.86227e-05	1	36.0441	365.215	
train_model_4a7e5_00010	TERMINATED	127.0.0.1:4960	0.001	0.0158172	8	8	0.000962394	1	57.6775	383.176	
train_model_4a7e5_00011	TERMINATED	127.0.0.1:4960	0.0001	0.0741538	8	5	0.000744995	1	6.58825	384.522	
train_model_4a7e5_00012	TERMINATED	127.0.0.1:4960	0.01	0.062308	24	6	0.000413833	1	39.5412	356.298	
train_model_4a7e5_00013	TERMINATED	127.0.0.1:4960	0.001	0.0732991	24	6	0.000822452	1	44.9731	380.849	
train_model_4a7e5_00014	TERMINATED	127.0.0.1:4960	0.0001	0.083812	24	3	0.000295798	1	4.36821	384.008	
train_model_4a7e5_00015	TERMINATED	127.0.0.1:4960	0.01	0.0126668	32	7	0.000970163	1	46.4767	356.458	
train_model_4a7e5_00016	TERMINATED	127.0.0.1:4960	0.001	0.0467678	32	2	2.01279e-05	1	18.7118	383.053	
train_model_4a7e5_00017	TERMINATED	127.0.0.1:4960	0.0001	0.0486875	32	7	0.000602337	1	43.2943	384.209	
train_model_4a7e5_00018	TERMINATED	127.0.0.1:4960	0.01	0.0223889	8	7	0.000438683	1	49.3722	354.329	

Test data can be split from data to measure accuracy at a higher prediction than validation.

Results:

Results were measure in 3 ways:

- RMSE value - the lower the more accurate the predictions were
- Ttest t value - the lower the value is the lower the difference between the teacher and

$$T = \frac{mean1 - mean2}{\frac{s(diff)}{\sqrt{(n)}}}$$

the predictions were

- Ttest p-value - the confidence level of whether the two groups are different. The closer to 1 the more accurate the predictions were.

Results:

```
The CV score is 110.42551
RMSE ttest result: 0.006068847438167256
RMSE ttest pvalue: 0.995157953331592
```

Tuning results:

```
best config: {'lr': 0.01, 'momentum': 0.06869814311075284, 'n_d': 8,
'weight_decay': 9.390065766228706e-05, 'n_steps': 2}
```


Combined CV & BI:

This model is similar to BI model, but uses a features matrix from the Rishon branch only and uses additional features from CV data.

The features added from CV data are:

- Percent actions of engagement with the product out of all actions.
- Dayin week
- Age of costumer
- Mean location (based on x,y) of actions which engage with the product
- Mean location (based on x,y) of actions which don't engage with the product

Date from CV data and BI data matched per row of feature matrix.

Tabnet hyperparameters were reoptimized using ray tuning for this model.

Results:

```
The CV score is 334.95236
RMSE ttest result: -0.5382598877087904
RMSE ttest pvalue: 0.5910276308855287
```

Tuning results:

```
best config: {'lr': 0.01, 'momentum': 0.017508676653025766, 'n_d': 32,
'weight_decay': 9.73785653440169e-05, 'n_steps': 9}
```

Total runtime:

```
Runtime 4579.555155038834
```

Since the data is much smaller we see a much less accurate predicting.

Appendix:

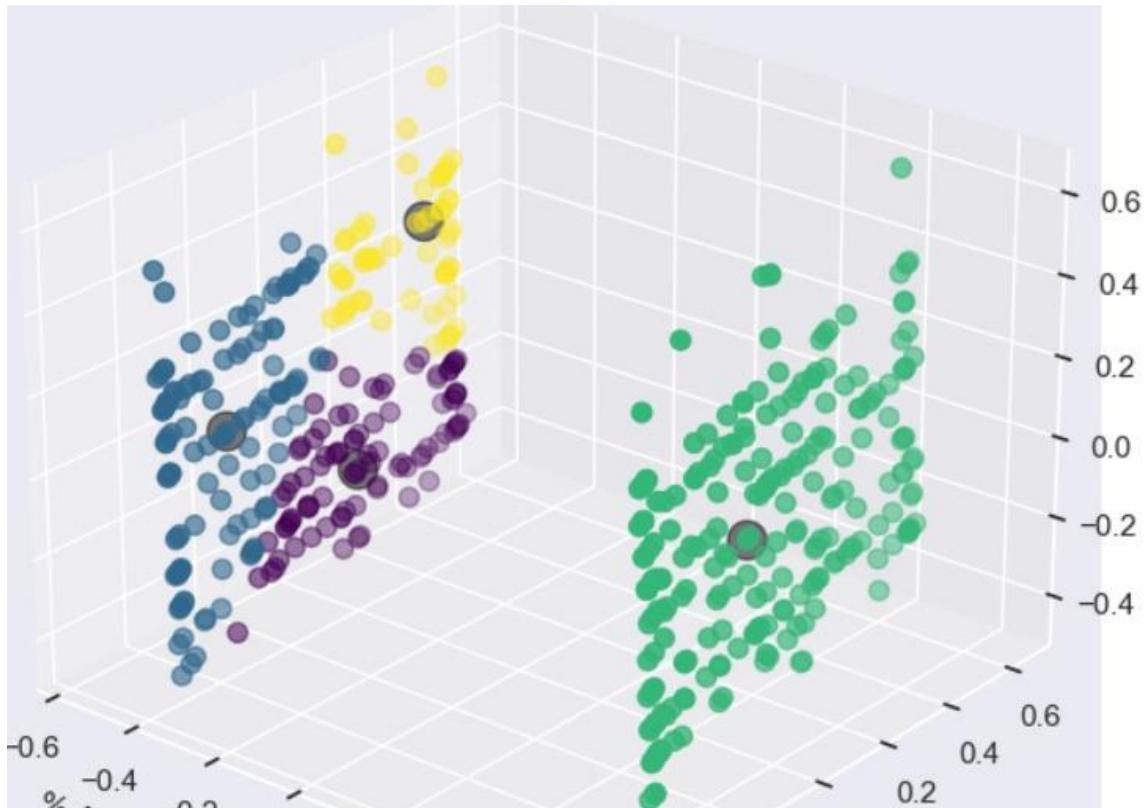
Epochs in tuning and running were reduced due to high run time. Multiple GPUs and bigger VRAM would reduce runtime and allow it to increase..

CV data:

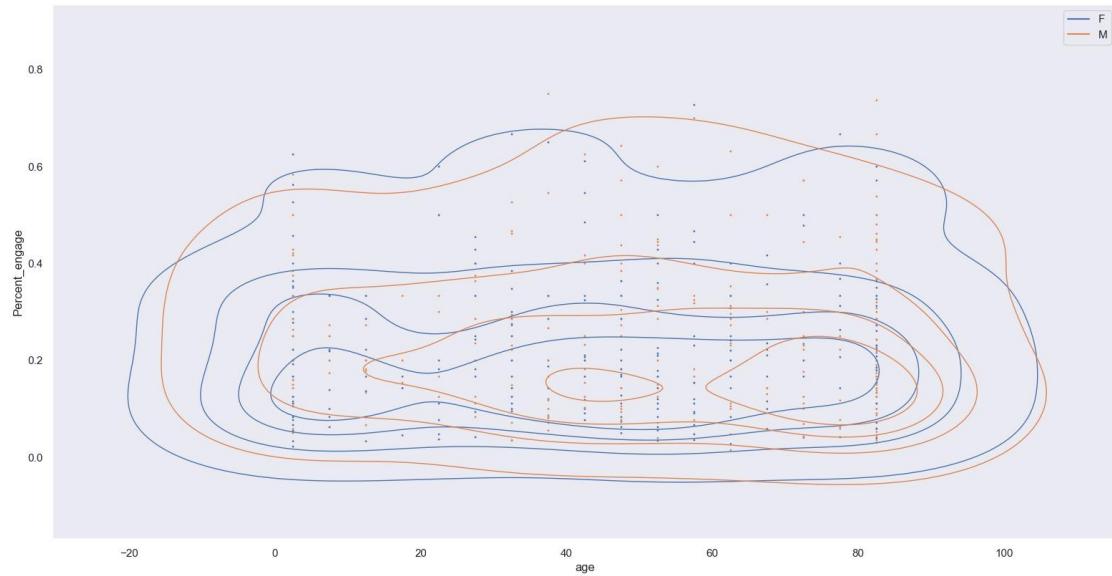
If data was a continuous procedure of shopping with consistent scoring per item purchase. Ideally I would use the “actions_process” function. This function returns the actions the customer did until they put the item in the cart. By looking at the number of actions the consumer did until they chose the item the model would be able to detect behavior that leads to item purchase. Binary returns an array of each costumer how many engagements with products they did before put to cart. While multiclass returns a list of all the actions done before put to cart.

Other clustering methods that were considered:


- PCA - in order to use all CV data and reduce dimension. This was not chosen due to difficulty to explain results.



- No clustering visualization - we can see that there is age and gender difference in engagement.



- Highest feature correlation to Pidyun:



BI features:

- I started from looking at two weeks before and after.

```
timeframe=(date-datetime.timedelta(days=delta_days)).strftime("%Y-%m-%d") + ' ' +
((date+datetime.timedelta(days=delta_days)).strftime("%Y-%m-%d"))
```

Afterwards I realized the future cannot affect people's shopping habits in the present so I only looked at 14 days in the past. The variable chosen is the mean slope of google trends of a word 3 days prior to the date.

- Other external features ideas: Traffic, special event (e.g sports, terror ect.), construction around the branch or in it.

General:

To save run time only "משומר בשר/דגים" category was run. To change to all data delete line:

```
df = df[df.category == category_get_only] # Get only specific category
```

- the candidate needs to *choice* - needs to be corrected to choose*.
- BI data from sales (15 excel files in Drive each represents a *signle* day). - needs to be corrected to single*
- Using database (SQL/NoSQL) is a bunos -needs to be corrected to bonus*
- Excels are long and big and take a long time to run - recommended to work with .csv